

Under the hood: System Calls

The term “function” is used many places in this book and needs no additional explanation here. The term “system call” has also been used in many places, often with a comment that making a system call through the C runtime is actually just calling an ordinary function acting as a wrapper. An explanation of this has been promised...

Wrappers

“Wrapper” is a term used to describe a function which hides the details of something else, often another function or functions. Hiding details is a form of abstraction and can be a good thing. Broadly speaking, an API (Application Programmer’s Interface) is itself another example of wrappers in common use.

C runtime as a wrapper

Many C runtime functions are just wrappers for system calls. For example if you call `open()` from the C runtime, the function will perform a few bookkeeping operations and then make the actual system call.

What IS a system call?

The short answer is a system call is a sort-of function call that is serviced by the operating system itself, within its own private region of memory and with access to internal features and data structures.

Our programs run in “userland”. The technical name for userland on the ARM64 processor is EL0 (Exception Level 0).

We can operate within the kernel’s space only through carefully controlled mechanisms - such as system calls. The technical name for where the kernel (or system) generally operates is called EL1.

There are two higher Exception Levels (EL2 and EL3) which are beyond the scope of this book.

Mechanism of making a system call

First, like any function call, parameters need to be set up. The first parameter goes in the first register, etc.

Second, a number associated with the specific system call we wish to make is loaded in a specific register (`w8`).

Finally, a special instruction `svc` causes a trap which elevates us out of userland into kernel space. Said differently, `svc` causes a transition from EL0 to EL1. There, various checks are done and the actual code for the system call is run.

A description of returning from a system call is beyond the scope of this book. Hint: just as there's a special instruction that escalates from EL0 to EL1, there is a special instruction that does the reverse.

What is the number associated with a particular system call?

Hard question.

In a perfect world, each Linux distribution would use the same set of system call numbers. But no.

This is the most comprehensive list of system call numbers we have seen. It shows system call numbers for many architectures and distributions.

Example: calling `getpid()`

The system call `getpid()` fetches the running process's process ID. Every executing entity has one.

We present four different versions of the same program:

1. Written in C++
2. Written in C
3. Written using C runtime from assembly language
4. Calling the system call directly from assembly language

Written in C++

```
#include <iostream> // 1
#include <unistd.h> // 2
// 3
using std::cout; // 4
using std::endl; // 5
// 6
int main() { // 7
    cout << "Greetings from: " << getpid() << endl; // 8
    return 0; // 9
} // 10
```

Written in C

```
#include <stdio.h> // 1
#include <unistd.h> // 2
// 3
int main() { // 4
    printf("Greetings from: %d\n", getpid()); // 5
}
```

```

    return 0;           // 6
}                       // 7

```

Written in assembly language using C runtime

```

        .global main           // 1
        .text                  // 2
        .align 2               // 3
                                // 4
main:    stp    x29, x30, [sp, -16]! // 5
        bl     getpid          // 6
        mov    w1, w0          // 7
        ldr    x0, =fmt        // 8
        bl     printf          // 9
        ldp    x29, x30, [sp], 16 // 10
        mov    w0, wzr         // 11
        ret                                // 12
                                // 13
        .data                  // 14
fmt:     .asciz "Greetings from: %d\n" // 15
                                // 16
        .end                   // 17

```

And finally: calling the system call directly

```

        .global main           // 1
        .text                  // 2
        .align 2               // 3
                                // 4
main:    stp    x29, x30, [sp, -16]! // 5
        mov    x8, 172          // 6 // getpid on ARM64
        svc    0                // 7 // trap to EL1
        mov    w1, w0          // 8
        ldr    x0, =fmt        // 9
        bl     printf          // 10
        ldp    x29, x30, [sp], 16 // 11
        mov    w0, wzr         // 12
        ret                                // 13
                                // 14
        .data                  // 15
fmt:     .asciz "Greetings from: %d\n" // 16
                                // 17
        .end                   // 18

```

We chose `getpid()` because it doesn't require any parameters. Using the C runtime, we simply `bl` to it. Calling the system call directly is different in that

we must first load `x8` with the number that corresponds to `getpid()` for the AARCH64 architecture.

Consulting this awesome website, we find that the number we want is 172.

Linux system calls table for several architectures

- system calls listed in man 2 unimplemented are not listed
- there is now Python module for querying this data: [system-calls](#)

Author info

System calls table is maintained by [Marcin Juszkievicz](#) — AArch64/Arm developer working at Red Hat as Linaro assignee.

How to help

Sources used to generate table are available in [git repository at github](#). Patches are always welcomed.

If this table helped you in some project then please leave a comment in [issue on GitHub](#).

Build info

Table generated on **27 November 2022 08:52 UTC** using data from **6.1.0-rc6** kernel source.

disable architectures ▾ Search:

system call	arm64	arm	armoabi	x86_64	x32	i386	powerpc
getpid	172	20	9437204	39	1073741863	20	20

Figure 1: here

The constant specific to the system call we want is loaded into `x8`. Recall that `x0` through `x7` are scratch registers.

Then on line 7, the `svc` with the argument 0 initiates the escalation from EL0 to EL1 where the kernel implements our desired functionality and returns to us.

Review

System calls are functions implemented inside the operating system.

To get there, at some point perhaps behind a wrapper function found in the CRT (C Run Time library), a distro specific system call number is placed in `x8` with other scratch registers getting the system call's documented parameters and the `svc` instruction is executed with argument 0.

Preference

We suggest using the CRT wrapper functions where possible because:

- They are easier to code
- They are portable between distributions of the OS