

Section 2 / Floating Point Literals

Recall that all AARCH64 instructions are 4 bytes long. Recall also that this means that there are constraints on what can be specified as a literal since the literal must be encoded into the 4 byte instruction. If the literal is too large, an assembler error will result.

Given that floating point values are always at least 4 bytes long themselves, using floating point literals is extremely constrained. For example:

```
fmov    d0, 1      // 1
fmov    d0, 1.1    // 2
```

Line 1 will pass muster but Line 2 will cause an error.

To load a `float`, you could translate the value to binary and do as the following:

```
        .text                                // 1
        .global main                        // 2
        .align 2                            // 3
main:    str    x30, [sp, -16]!              // 4
        ldr    s0, =0x3fc00000             // 5
        fcvtn  d0, s0                      // 6
        ldr    x0, =fmt                    // 7
        bl     printf                      // 8
        ldr    x30, [sp], 16               // 9
        mov    w0, wzr                     // 10
        ret                                // 11
        .data                                // 12
fmt:     .asciz  "%f\n"                     // 13
        .end                                // 14
        .end                                // 15
        .end                                // 16
```

The above code is kind of found here - the file is used for miscellaneous testing.

Line 6 puts the translated value of 1.5 into `s0` (since we are thinking of the value as a `float` it goes in an `s` register). The assembler performs some magic getting a 32 bit value seemingly fit into a 32 bit instruction. See below.

Line 7 converts the single precision number into a double precision number for printing.

printf() only knows how to print double precision values. When you specify a float, it will convert it to a double before emitting it.

Translating `floats` and `doubles` by hand isn't a common practice for humans, though compilers are happy to do so.

Instead for us humans, the assembler directives `.float` and `.double` are used more frequently to specify `float` and `double` values putting them into RAM.

The following example prints an array of floats and doubles:

```

        .global main                                // 1
        .text                                       // 2
        .align 2                                    // 3
                                                // 4
counter .req    x20                                // 5
dptr    .req    x21                                // 6
fptr     .req    x22                                // 7
        .equ    max, 4                             // 8
                                                // 9
main:    stp     counter, x30, [sp, -16]!           // 10
        stp     dptr, fptr, [sp, -16]!           // 11
        ldr     dptr, =d                           // 12
        ldr     fptr, =f                           // 13
        mov     counter, xzr                       // 14
                                                // 15
1:        cmp     counter, max                      // 16
        beq     2f                                  // 17
                                                // 18
        ldr     d0, [dptr, counter, lsl 3]        // 19
        ldr     s1, [fptr, counter, lsl 2]        // 20
        fcvt    d1, s1                             // 21
        ldr     x0, =fmt                           // 22
        add     counter, counter, 1                // 23
        mov     x1, counter                        // 24
        bl     printf                              // 25
        b       1b                                  // 26
                                                // 27
2:        ldp     dptr, fptr, [sp], 16             // 28
        ldp     counter, x30, [sp], 16             // 29
        mov     w0, wzr                             // 30
        ret                                         // 31
                                                // 32
        .data                                       // 33
fmt:     .asciz  "%d %f %f\n"                      // 34
d:       .double 1.111111, 2.222222, 3.333333, 4.444444 // 35
f:       .float  1.111111, 2.222222, 3.333333, 4.444444 // 36
                                                // 37
        .end                                         // 38

```

The above code is found here.

A number of interesting things in this source code:

- We use `.req` to give symbolic names to various registers. This can help you in remembering which register is being used for what purpose.

- We use `.equ` to encode a small integer literal value to give it a symbolic name, eliminating the use of a “magic number.”
- Lines 19 and 20 use address arithmetic to march through an array of doubles (8 bytes each) and an array of floats (4 bytes each).

Line 19 is equivalent to:

```
// ldr    d0, [dptr, counter, lsl 3]
d0 = dptr[counter];
```

`counter` is multiplied by 8 then added to `dptr`.

Line 20 is equivalent to:

`counter` is multiplied by 4 then added to `fptra`.

```
// ldr    s1, [fptra, counter, lsl 2]
s1 = fptra[counter];
```

Cool huh?

On Linux, just as `w/x0` through `w/x7` are scratch registers and used to pass parameters, `s/d0` and `s/d7` are as well beginning with the 0 register.

Fitting 32 bits into a 32 bit bag

This section is currently LINUX-centric - in the future it will address both native Apple and Linux equally.*

AARCH64 instructions are 32 bits in width. Yet, line 6 from this program reads:

```
ldr    s0, =0x3fc00000 // 6
```

This appears to show a 32 bit constant being held in an instruction that itself is 32 bits wide. Well, the Assembler does some magic. Let’s see what that magic is.

Build the program with the `-g` option to enable debugging using GDB.

```
% gcc -g t.s
```

Then launch GDB on the executable:

```
% gdb a.out
```

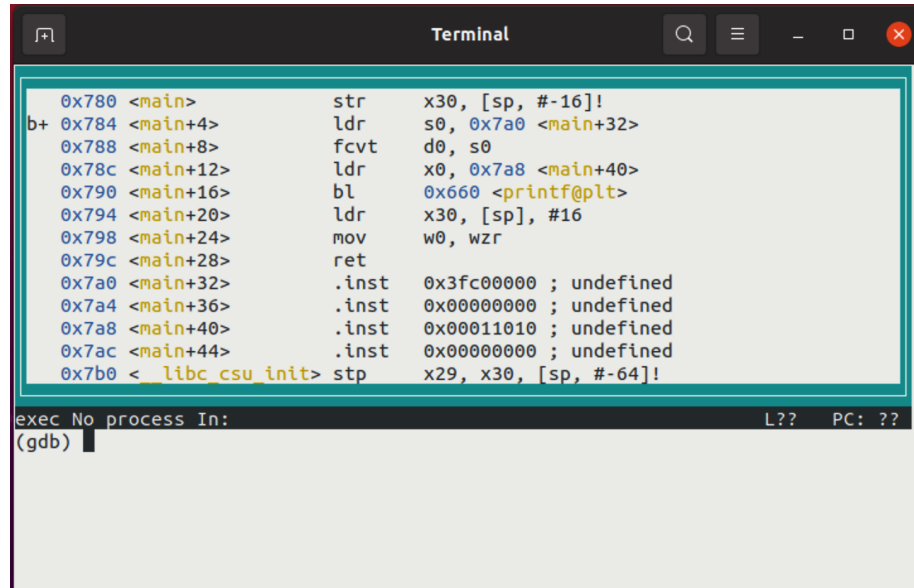
Set a breakpoint on line 6.

```
(gdb) b 6
Breakpoint 1 at 0x784: file t.s, line 6.
(gdb)
```

Enter a cool GDB layout (one of several cool layouts):

```
layout asm
```

You should see something like this:



```
0x780 <main>      str     x30, [sp, #-16]!
b+ 0x784 <main+4>   ldr     s0, 0x7a0 <main+32>
0x788 <main+8>     fcvt    d0, s0
0x78c <main+12>    ldr     x0, 0x7a8 <main+40>
0x790 <main+16>    bl      0x660 <printf@plt>
0x794 <main+20>    ldr     x30, [sp], #16
0x798 <main+24>    mov     w0, wzr
0x79c <main+28>    ret
0x7a0 <main+32>    .inst   0x3fc00000 ; undefined
0x7a4 <main+36>    .inst   0x00000000 ; undefined
0x7a8 <main+40>    .inst   0x00011010 ; undefined
0x7ac <main+44>    .inst   0x00000000 ; undefined
0x7b0 <__libc_csu_init> stp    x29, x30, [sp, #-64]!
```

exec No process in: L?? PC: ??
(gdb)

Figure 1: gdb01

We expected line 6 to read:

```
ldr     s0, =0x3fc00000
```

Instead we find:

```
b+ 0x784 <main+4>      ldr     s0, 0x7a0 <main+32>
```

Scan downward to find 0x7a0:

```
0x7a0 <main+32>      .inst   0x3fc00000 ; undefined
```

Hey look! Here's our literal float. The `.inst` is an ARM specific GNU assembler directive says: `^_(-)_/^`.

Note, the encoded “instruction” does not have to make any sense - instead the compiler has emitted a make believe instruction that happens to have the value of our literal.

What we're seeing the actual line 6 doing is reaching ahead a short distance to load the value of another location in memory where our constant is really found.

Let us take this explanation further. Notice we see:

```
0x78c <main+12>      ldr     x0, 0x7a8 <main+40>
```

where we expected:

```
ldr      x0, =fmt
```

Scan down to 0x7a8:

```
0x7a8 <main+40>      .inst  0x00011010 ; undefined
```

x0 is serving as a pointer to the format string of a call to `printf()`. Let's follow the pointer...

```
(gdb) x/s 0x00011010
0x11010:      "%f\n"
(gdb)
```

Magic.